

Introduction to Artificial Neural Networks

Introduction based on the Bachelor Thesis of Marius Felix Killinger (2014)

Contents

1	Artificial Neural Networks	4
1.1	Origin and Motivation	4
1.2	Biological Background	4
1.3	Neural Networks as a Mathematical Model	5
1.4	Application to Machine Learning	7
1.5	Training	9
1.5.1	Backpropagation	10
1.5.2	Optimisation	12
1.6	Weight Initialisation	15
2	Convolutional Neural Networks	17
2.1	Origin and Motivation	17
2.2	Network Architecture	18
2.2.1	Convolution	18
2.2.2	Max-Pooling	19
2.3	Operation of a CNN	19
2.3.1	Providing training data to a CNN for image-to-image tasks	21

Nomenclature

Symbol	Description
NN	Neural Network
CNN	Convolutional Neural Network
MSD	Mean Squared Deviation
SGD	Stochastic Gradient Descent
\mathbf{x}	Input vector
$\mathbf{x}^{(n)}$	n -th data example, input to network
$\mathbf{x}_{(i)}$	input of i -th layer
\mathbf{y}	Output activation vector
$\mathbf{y}^{*(n)}$	Target label of n -th data example

\mathbf{W}	Weight matrix of connections between two layers (can implicitly include biases)
\mathbf{b}	Explicit bias vector
\mathbb{W}	Set of all weights (and biases), set of parameters
\mathbf{a}	Vector of weighted sums of incoming signals: $\mathbf{W}^T \mathbf{x}$
$\mathbf{a}_{(i)}$	Vector of weighted sums of incoming signals of layer i
\mathbf{h}	Activation function to be applied to \mathbf{a}
\mathcal{E}	Loss/Error function
η	Learning rate, step size

Chapter 1

Artificial Neural Networks

1.1 Origin and Motivation

Many perspectives can be assumed for analysing Neural Networks (NNs), ranging from an biological perspective to a purely mathematical point of view. I will start to describe NNs from biological foundations and move on to their mathematical properties and applications for machine learning.

1.2 Biological Background

An important step towards understanding neural information processing was the electrophysiological analysis of the *squid giant axon* by [Young, 1938]. Subsequently a primitive neuron model was formulated by [McCulloch and Pitts, 1943], a comprehensive quantitative model of the Neuron has been proposed later by [Hodgkin and Huxley, 1952] which opened the path to model neural networks. It was known, that biological neurons are interconnected by synapses through which electrical signals are transferred. Depending on the incoming signal, neurons become active i.e. create new *spike*-signals which in turn stimulate other neurons' activities. The spiking neuron is an, at least 4-dimensional, highly non-linear dynamical system. Thus the combination of even a small number of such neurons to form a network by making synaptic connections, quickly results in very complex systems. Depending on the model parameters, this system can easily drift into the realm of chaos; purposeful computations are only possible in a small and sensitive corridor of parameters ([Lazar et al., 2009].

To reduce complexity and cope with computational limitations dramatical simplifications were made. A commonly used neuron model reduces the spiking neuron to an element that sums its input and applies a non-linear *activation-function* h . The activation of the model neuron can be interpreted as the mean firing rate of a biological neuron, neglecting individual spikes and their respective phases. The *activation* y , given the inputs

of n synaptic connections x_i and a weight w_i for each connection, is given by:

$$y = h\left(\sum_{i=1}^n w_i x_i\right) = h(a) \quad (1.1)$$

a is introduced as a shorthand for the weighted sum. For m neurons the weights are given by a $m \times n$ Matrix \mathbf{W} and the calculation is noted as follows:

$$\mathbf{y} = h(\mathbf{W}^T \mathbf{x}) = h(\mathbf{a}) \quad (1.2)$$

Where h denotes the activation-function applied component-wise. McCulloch-Pitts Neurons use a heaviside step function as activation-function. The binary output results in boolean logic networks. This turned out to be too restrictive and it was later replaced by continuous sigmoidal activation-functions (logistic function, hyperbolic tangent and rectified linear unit - the latter is not sigmoid in a stric sense). Additionally an affine offset called *bias* is added to the sum before applying the activation-function ($\mathbf{W}^T \mathbf{x} \rightarrow \mathbf{W}^T \mathbf{x} + \mathbf{b}$). The biases can be subsumed into the weight matrix by introducing an additional connection from an additional neuron for each layer that is always active (activity= 1). Although computer experiments are not implemented this way, it is convenient for the derivation of the Backprop rule, therefore I will neglect biases henceforth.




Figure 1.1: A single model neuron

1.3 Neural Networks as a Mathematical Model

For completeness it should be mentioned that a great deal of research in the field of neuroscience does not aim at making artificial neural networks abstract and simple, but instead tries to simulate the processes as close as possible to the biological original. However, this is not the goal in machine learning therefore I will set neurophysiological realism aside from now on.

Despite being also of great interest as mathematical models, often artificial NNs exclude *recurrent* synaptic connections and restricts networks to be acyclic. Such networks are called *feed-forward* and the output activations, given some input activations, can directly be calculated: the network is simply a way of specifying a ordinary mathematical *function*. In contrast, recurrent networks specify a dynamical system: determining the output requires solving coupled differential equations - or in a relaxed setting discrete iterations of value updates until some stopping criterion has been fulfilled; effectively this makes recurrent networks feed-forward networks again by unrolling the time dimension and truncating it. It is clear that recurrent neural networks have greater computational power but are also more difficult to use.

Feed-forward networks allow neurons to be arranged in *layers*: neurons in the same layer only receive their input from a common input vector and return an output vector. There is no interaction of neurons within a layer, they operate in parallel, thus the network is acyclic. The output of such a network can be calculated by nesting the $\mathbf{W}^T \mathbf{x}$ -multiplications and applying the activation-function h component-wise. In contrast to locally connected *convolutional* layers, such layers are called *fully-connected* layers. For d layers the *forward pass* is computed as:

$$\mathbf{y} = h(\mathbf{W}_{d-1 \rightarrow d}^T \dots h(\mathbf{W}_{1 \rightarrow 2}^T h(\mathbf{W}_{0 \rightarrow 1}^T \mathbf{x}))) \quad (1.3)$$

The indices of the weight matrices denote from which layer to which the synaptic connections project. The number of neurons in each layer may be arbitrary; the dimensions of the weight matrices change accordingly. The vector \mathbf{x} constitutes the input for layer 1 and \mathbf{y} is the output of the last layer d ; all layers in between are called hidden-layers. Input is not be counted as a separate layer, e.g. a one layer network consists of its input vector, the weight multiplication and the activation vector as output. All weight matrices (and bias vectors) together form the set of parameters \mathbb{W} of a NN.

We see from equation 1.3 that in the case of a linear activation-function (i.e. the identity) a neural network is merely a nested matrix multiplication. Then it can be simplified by a single matrix with rank lower or equal to the lowest rank of the individual weight matrices. Training such a network is comparable to PCA, although computationally inefficient. Likewise for non-linear activation-functions NN perform a non-linear PCA, now the computational effort can be justified by the increased expressive power of the model.

The first use of NNs in a pure machine learning context was the so called Perceptron by [Rosenblatt, 1958]. The Perceptron consists of a single hidden layer with a linear activation-function. In the most simple case it has only one output neuron. This can be used for a binary classification of multidimensional input. Because it is linear and has just one hidden layer, this classifier can reach optimal discrimination on linearly separable data only. However convergence is guaranteed in this case. The Perceptron is closely related to

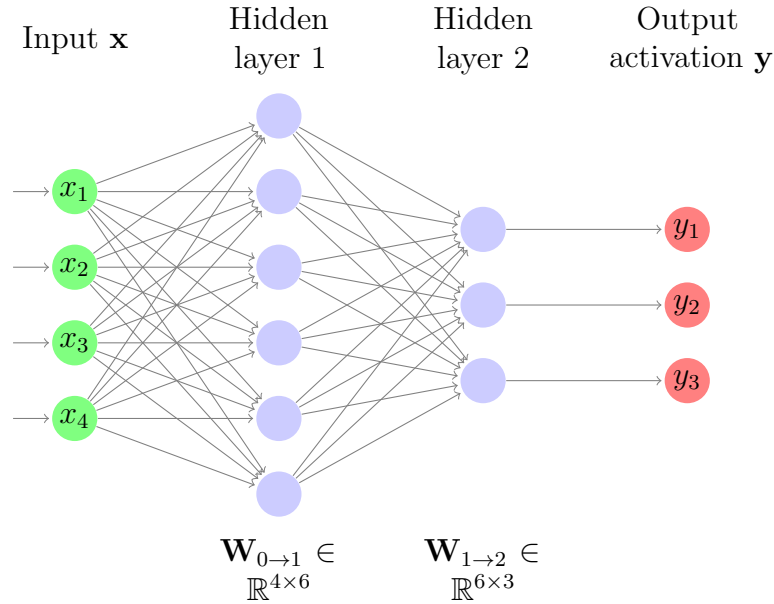


Figure 1.2: A 2-layer Neural Network

logistic regression and linear discriminant analysis.

It has been shown by [Cybenko, 1989] that a network with two hidden layers and a sigmoid activation in the first and linear in the second layer can approximate continuous functions on a compact subset of \mathbb{R}^m to an arbitrary precision (the required width of the layer may scale as bad as $O(\exp(m))$, in contrast deeper networks with more than two layers are much more effective in terms of required computations and model parameters). It is this property - universal function approximation - which makes NNs useful for many applications.

1.4 Application to Machine Learning

The most general setting in machine learning is a data set containing N examples. Each example is represented by a j -dimensional *feature*-vector \mathbf{x} . A feature can be any property that can be mapped to real or discrete values, e.g. for an image the features can be simply the pixel gray values. Each example can optionally be annotated with one or more labels \mathbf{y}^* which can be continuous or discrete. For all machine learning tasks the common rationale is to *generalise* information gathered on one subset of the data, the *training set*, and use the *learnt* information to perform well on new data which comes from a held out *test set*¹. The criterion for a model's quality is determined through the performance on

¹In a real world application the test set is only used to assess a model. When a model is accepted the data to be classified is genuinely new.

new data.

A variety of common tasks are:

Classification:

For a data example the membership of one (out of several) classes must be inferred from the j -dimensional feature-vector. E.g. predicting whether a person's income is above or below a certain value by providing demographic facts about a person.

Multi-label Classification:

The same as above, but here several categories (each of which comprise several possible classes) must be predicted. E.g. recognising an object's category and its colour independently.

Regression:

Instead of discrete class membership a real value must be predicted. E.g. the market value of a house given facts about the location and house.

Reconstruction:

Disrupted data must be restored. E.g. noise reduction in images.

Prediction:

Similar to Regression but in a temporal context. E.g. predicting the air temperature of the next day given data of the previous day's environmental conditions.

Compression / Dimensionality Reduction:

The dimensionality of the feature space must be reduced under the constraint that the original features can be reconstructed well from the compressed representation or that certain information of interest is retained.

System Control:

For given sensor input the best reaction of actuators in order to maintain specified process parameters. E.g. controlling pressure and temperature in a chemical reactor.

All of these tasks are linked closely: reconstructing a disrupted image resembles restoring a compressed version (both require good models for image content), to control a system it is necessary to make predictions about the system parameter's future evolution and so on. The unifying property of these tasks is that they can be accomplished by mapping input features to a set of output variables. Due to their function approximation capabilities NNs offer a unifying framework to deploy this mapping function.

The common sigmoid functions have the co-domain $[-1, 1]$ or $[0, 1]$ and are virtually

flat for large absolute arguments. Hence it is appropriate to normalise input features². For categorical features two options exist:

1. Using one neuron per category and divide its value range into "bins" corresponding to the different classes. This is applicable if classes have an order relation (ordinal data).
2. For k possible classes k neurons are used: only the neuron corresponding to the correct class receives input 1 and all others 0 (One-of- k -coding). This is always possible, in particular also for nominal data.

For the output, in general the same applies; however, one thing must be taken into account: it is desirable (e.g. for comparison across different classifiers) to have a probability distributions over the different classes. The class with the highest probability wins and the example is predicted to have the corresponding label. However, the output of the network is in general neither non-negative nor normalised. This can be achieved by applying a *softmax*-transform to the k output neurons of a class instead of a sigmoid activation-function:

$$\mathbf{y} = \frac{e^{\mathbf{x}}}{\sum_{i=1}^k e^{x_i}} \quad \text{and} \quad \underset{i}{\operatorname{argmax}} y_i \quad \text{the most probable class} \quad (1.4)$$

\mathbf{x}_n is the vector of activations of the last layer. In the numerator the exponentiation is applied component-wise. The result \mathbf{y} is an empirical categorical probability distribution over the class memberships. Taking the argmax presumes the classes are encoded by the vector's positional index.

1.5 Training

The aim of the training is to find a good parameter configuration so that the network maps the training examples in the desired way to output. This target is formulated by minimising an *error-function* $\mathcal{E}(\mathbb{W})$ which is a measure of discrepancy between actual and desired network output. A standard error-function is the mean squared distance error (MSD) averaged over all training examples:

$$\mathcal{E}(\mathbb{W}) = \frac{1}{2} \frac{1}{N} \sum_{n=1}^N \|\mathbf{y}(\mathbf{x}^{(n)}; \mathbb{W}) - \mathbf{y}^{*(n)}\|^2 \quad (1.5)$$

²This is also necessary in order to use initialisation schemes for the weights which assume the variance is in the order of 1.

When using the softmax to yield a probabilistic output, the true label of example n can be regarded as a probability distribution Q_n . Then it is possible to use the cross entropy between the two distributions. This has also the benefit to allow for probabilistic labelling. For a single category with k possible classes \mathbf{y} is a length- k vector and we write:

$$\mathcal{E}(\mathbb{W}) = -\frac{1}{N} \sum_{n=1}^N \sum_{i=1}^k \log(P_n(Y = i | \mathbf{x}^{(n)}; \mathbb{W})) Q_n(Y = i) \quad \text{with:} \quad (1.6)$$

$$P_n(Y = i | \mathbf{x}^{(n)}, \mathbb{W}) = \mathbf{y}_i(\mathbf{x}^{(n)}; \mathbb{W}) \quad (1.7)$$

$$Q_n(Y = i) = \delta_{i, y^*(n)} \quad (1.8)$$

For multi-category classification this can be generalised by doing the same for each category and summing or averaging the errors³.

Equation 1.8 tells that for the label distribution all the mass is concentrated on the true class and others have 0 probability. Because of this, the cross entropy is identical to the Kullback-Leibler divergence and moreover we can simplify the error-function:

$$\mathcal{E}(\mathbb{W}) = -\frac{1}{N} \sum_{n=1}^N \log(\mathbf{y}_{(y_i^*(n))}(\mathbf{x}^{(n)}; \mathbb{W})) \quad (1.9)$$

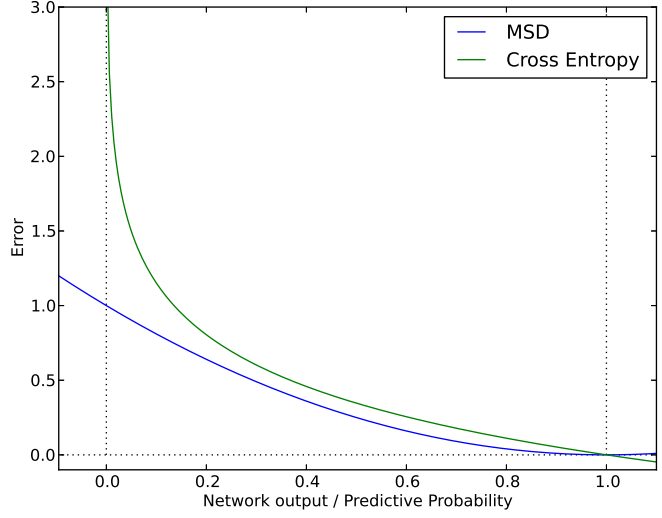
Apparently the MSD has unlimited range and can thus be used to train the network for real-valued regression tasks. The cross entropy is only defined for values from $[0, 1]$, making it applicable when the target is a (class) probability. As depicted in figure 1.3 the cross entropy grows beyond limits if the predictive probability of the true class approaches zero. This is an advantage because it gives a much stronger correction signal than the MSD in order to update the parameters such that the probability of the true class increases. The MSD is still of finite value which makes the learning less radical.

1.5.1 Backpropagation

After defining an appropriate error-function the gradient must be computed to minimise it. *Backpropagation* is a computationally efficient universal scheme for propagating the errors like messages sent from the output neurons back through the network to obtain the gradient. It was first applied to NNs by [Rumelhart et al., 1986].

³Constants in front of the error-functions can be chosen arbitrarily, their effect is to scale the gradient, but they do not change the direction.

Figure 1.3: Comparison of the MSD and the cross entropy error-functions. This plot shows the error versus the network output for the true class, which is desired to be 1. The cross entropy enforces a more rigid training to stay away from low probabilities for the true class. Note that this difference is not due to scaling of the functions, instead it is caused by the singularity at 0 of the logarithm. Besides that, the MSD becomes flatter near the target value which additionally makes training slower in comparison.



Backprop can be seen as a recursive algorithm that does one recursion per layer. The derivation starts at the output of a single-layer NN. Because differentiation is linear, it is sufficient to consider only a single term of the sum-like error-function for the derivation; the generalisation to all terms of the sum is clear. Biases can be included in the weight matrices as mentioned earlier. The entry point of the recursion is the differentiation of the error w.r.t the activation of the output neuron layer:

$$\nabla \mathcal{E} \Big|_{\mathbf{w}} = \frac{\partial \mathcal{E}(\mathbf{W})}{\partial \mathbf{W}} = \frac{\partial \mathcal{E}(\mathbf{W})}{\partial \mathbf{y}(\mathbf{x}; \mathbf{W})} \frac{\partial \mathbf{y}(\mathbf{x}; \mathbf{W})}{\partial \mathbf{W}} \quad (1.10)$$

$$= \frac{\partial \mathcal{E}(\mathbf{W})}{\partial \mathbf{y}(\mathbf{x}; \mathbf{W})} \frac{\partial h(\mathbf{a})}{\partial \mathbf{a}} \frac{\partial \mathbf{a}}{\partial \mathbf{W}} \quad \mathbf{y} \text{ is } \mathbf{a} \text{ with the non-linearity applied} \quad (1.11)$$

$$= \frac{\partial \mathcal{E}(\mathbf{W})}{\partial \mathbf{y}(\mathbf{x}; \mathbf{W})} \frac{\partial h(\mathbf{a})}{\partial \mathbf{a}} \mathbf{x}^T \quad \mathbf{a} \text{ is linear in } \mathbf{W} \quad (1.12)$$

The first factor is determined by the true labels and our choice of error-function. The remainder is the derivative of the activation w.r.t the weights. Note that \mathbf{x} is transposed and all products are general matrix multiplications⁴. Consequently the gradient always has the same dimension as the weights w.r.t. which it was calculated. It is efficient to use an activation-function that has an easy to evaluate derivative. For MSD-error and logistic activation $(1 + \exp(-a))^{-1}$ this is:

$$\frac{\partial \mathcal{E}(\mathbf{W})}{\partial \mathbf{W}} = (\mathbf{y} - \mathbf{y}^*) h(\mathbf{a}) (1 - h(\mathbf{a})) \mathbf{x}^T \quad (1.13)$$

⁴In this case this is the outer product, which maps two vectors to a matrix.

To generalise the calculation to d layers we introduce an abbreviation for *errors* ϵ_i for layer i : the gradient of the error-function w.r.t to the input of layer i $\mathbf{x}_{(i)}$. The error of the output layer is the partial derivative of the error-function \mathcal{E} w.r.t. to the output layer's activation \mathbf{y} . The gradient of the next layer must be, by chain rule, multiplied with the "previous" error:

$$\epsilon_d = \frac{\partial \mathcal{E}(\mathbb{W})}{\partial \mathbf{y}(\mathbf{x}; \mathbb{W})} \quad (1.14)$$

$$\epsilon_i = \frac{\partial \mathcal{E}(\mathbb{W})}{\partial \mathbf{x}_{(i)}} = \epsilon_{i+1} \frac{\partial h(\mathbf{a}_{(i)})}{\partial \mathbf{a}_{(i)}} \mathbf{W}_{i \rightarrow i+1} \quad (1.15)$$

$$\frac{\partial \mathcal{E}(\mathbb{W})}{\partial \mathbf{W}_{i \rightarrow i+1}} = \epsilon_{i+1} \frac{\partial h(\mathbf{a}_{(i)})}{\partial \mathbf{a}_{(i)}} \mathbf{x}_{(i)}^T \quad (1.16)$$

The first equation gives us the initial value and the second one a recursion recipe. From the errors the gradients w.r.t to the weights can be found in the third step.

The full gradient w.r.t all parameters \mathbb{W} is the set of all the gradients of the individual layers. By consequently applying the chain rule and *reusing* the errors from the previous layer, Backprop enables us to determine the analytical gradient of a very complicated function (note that the activation-function can be arbitrary). The above derivation only looks at one single training example; by linearity of the error-function and differentiation the generalisation to an arbitrary number of examples is clear. A great benefit of this is the ability to parallelise the individual Backprops since there is no communication required between the training examples.

The cost of a Backprop-pass is $\mathcal{O}(\text{card}(\mathbb{W}))$ - the same as for a forward pass. This is very efficient as is illustrated by the following comparison: estimating the gradient from finite differences requires a perturbation in every parameter and evaluation how the error-function reacts. This yields $\mathcal{O}(\text{card}(\mathbb{W})^2)$ and is in addition numerically error-prone [Bishop et al., 2006, chap. 5.3.3].

1.5.2 Optimisation

Once the gradient is available, it is used to minimise the error. It should be noted that the error-function is non-convex and has a great number of local minima which makes optimisation hard. There also exist a great number of equivalent local minima for any given NN⁵: the weights exhibit symmetries such as exchanging pairs of neurons or flipping

⁵For a two-layer network with M hidden neurons, $2^M M!$ totally equivalent weight configurations exist [Bishop et al., 2006, chap. 5.1.1].

signs of incoming and outgoing weights at a neuron. This alone should make it clear that the concept of finding a unique maximum must be abandoned for NNs.


There exists a multitude of variants for optimisation. In general they can be subdivided into full-*batch* and mini-batch methods. The former computes the gradient using of the full training set (summing up all terms in the error-function) the latter only considers batches of a fixed number of training examples.

Given some initial weights at the start, the most basic scheme is a gradient descent optimisation with the following parameter updates:

$$\mathbb{W}_{t+1} = \mathbb{W}_t - \eta \nabla \mathcal{E} \Big|_{\mathbb{W}_t} \quad (1.17)$$

η is a meta parameter, the *learning rate*. The minus sign ensures that changes are made to minimise the error-function (for sufficiently small enough η , for too high η increases can be observed). This scheme can be used for both training modes; for full-batch it is known as standard Backpropagation.

Figure 1.4: Illustration of gradient descent on a 2-D error surface. For more dimensions the problem of multiple local optima is even more severe. [source: <http://blog.datumbox.com/tuning-the-learning-rate-in-gradient-descent/>]



In the following an overview of modifications to this scheme is given:

Stochastic Gradient Descent (SGD)

This scheme is most popular. Because mini-batches are much smaller (including batches of size one) the evaluation of error and gradient is significantly faster than for the full training set. The mini-batches are randomly sampled for each iteration step. This gives a noisy estimate of the gradient as the batch is not perfectly representative of the training set. However, the noise can help to escape a local minimum and jump to a new, better minimum (similar to the rationale behind simulated annealing). During the end of the training this noise might have a negative effect and disrupt the very fine tuning. To fix this the learning rate can be lowered gradually to make smaller changes at the end. By choosing the batch-size the noise versus speed

trade-off can be regulated⁶.

Momentum

To make the updates more robust against noise and help a more “direct” movement in parameter space a momentum term can be added:

$$\mathbb{W}_{t+1} = \mathbb{W}_t - \eta \left((1 - m) \nabla \mathcal{E} \Big|_{\mathbb{W}_t} + m \nabla \mathcal{E} \Big|_{\mathbb{W}_{t-1}} \right) \quad 0 \leq m < 1 \quad (1.18)$$

If the error-function is visualised as a height profile over the parameter space (as in figure 1.4), the trajectory of this scheme resembles a massive ball that rolls downhill but, by means of its momentum, is able to roll across little ridges, and will not change its direction abruptly.

Resilient Backprop (RPROP)

Setting the training rate in order to advance fast but also not to destroy what has already learnt, by going too far, is difficult. RPROP tries to fix this problem by estimating good learning rates for *each* parameter individually: if the gradient for one parameter has changed sign from the last to the next iteration, the learning rate is decreased, if the sign stays the same it is increased.

Second-order optimisation:

Newton Algorithm

A good estimate of the individual learning rates can be computed from the second derivative of the error, the Hessian \mathbf{H} :

$$\mathbb{W}_{t+1} = \mathbb{W}_t - \eta \left(\mathbf{H}^{-1} \nabla \mathcal{E} \right) \Big|_{\mathbb{W}_t} \quad (1.19)$$

This accounts for the fact that curvature is different in different dimensions of the parameter space. For strongly curved directions smaller steps are necessary, for more linear directions greater steps can be made. η is still needed to scale the overall progress. Unfortunately the computation of the Hessian and even more its inversion is very costly and does not scale well for deep networks.

Conjugate Gradient Descent (CG)

A general problem of simple descent methods is that they are likely to move along a zigzag-trajectory (because the Hessian is in general not diagonal). Directions which lie orthogonal w.r.t to the Hessian are called conjugate. There exist cheap heuristics for estimating conjugate directions by starting an iteration from a standard descent

⁶Note that computational cost increases linearly with batch-size whereas the variance decreases only as $O(\frac{1}{\sqrt{N}})$

direction (Polak-Ribiere coefficient). However, these heuristics are only effective if the update steps are close to optimal step size. Thus it is necessary to make a *line search* along the current direction, which has a large computational footprint. Besides, CG works best for full-batch learning. Whether mini-batches can be used depends on the particular task and also on how many examples are included in a mini-batch. Nonetheless, it is possible to compete with SGD, despite bigger batches and the line search, because the quality of the steps is much better. For CG, the learning rate must not be set explicitly.

Broyden-Fletcher-Goldfarb-Shanno method (BFGS)

This methods uses heuristics to approximate the inverse Hessian directly from the history of gradients. It is therefore much less complex than the Newton method, but in contrast to CG it requires storage of the approximated inverse Hessian which can become infeasible for large networks. It is intended for full batch learning.

This overview is not complete, many more heuristics and also certain unsupervised pre-training methods exist. For more details refer to [LeCun et al., 2012].

1.6 Weight Initialisation

Most commonly weights are initialised by random. Assuming input is normalised to have unit variance, it is desirable to keep that variance also for the output of the neurons of the first layer and so on for following layers. In this regime the sigmoid function is approximately linear. If inputs are too small (too small variance) the gradients are too weak, similarly for big inputs the sigmoids are saturated and also give small gradients. This notion gives rise to the picture that the network first must learn the easy linear mappings and non-linear components tune in at a later stages of the training. Assuming the N -dimensional input \mathbf{x} is uncorrelated and normalised, the standard deviation of the weighted sum a is:

$$\sigma_a = \left(\sum_{i=1}^n W_i \sigma_i \right)^{\frac{1}{2}} = \left(\sum_{i=1}^n W_i \right)^{\frac{1}{2}} \quad (1.20)$$

To make this unity, the distribution of weights should have the standard deviation:

$$\sigma_W = N^{-\frac{1}{2}} \quad (1.21)$$

However, there exist different heuristics, such as additionally taking into account the number of outgoing connections, or fixing the variance of the initial weights as a global constant. Strictly, all these considerations are depending on the input preparation and the

particular activation-function⁷.

⁷E.g. for rectified linear units (i.e. $h(x) = \max(0, x)$) the above reasoning is not applicable.

Chapter 2

Convolutional Neural Networks

2.1 Origin and Motivation

Convolutional Neural Networks (CNN) are a class of network topologies specifically designed for image processing. The input vector \mathbf{x} is reshaped into a matrix so that neuron activations in the input layer directly correspond to pixel intensities in the image; in deeper layers the neurons correspond to wider areas, not only single pixels. The key idea is to take advantage of the spatial ordering in images. CNNs can be viewed from two perspectives: first they provide a regularisation method by enforcing sparse connectivity, second they mimic the topology of the primary visual cortex which is by far the best known image processing unit. The biological point of view is also how they evolved historically:

[Hubel and Wiesel, 1962] discovered neurons in animals' visual cortex that respond to lines at different angles (simple cells) and neurons that respond to the activations of those simple cells, but with some tolerance as to where exactly the response came from. The latter results in a integration across a *field of view* (complex cells). [Fukushima, 1980] was inspired by these findings and designed a NN accordingly. His “Neocognitron” was later simplified by [LeCun et al., 1998].

For machine learning, one of the motivations was to restrict the number of model parameters: in common networks the number of parameters scales with the product of neuron count in two adjacent layers; this can quickly lead to a very complicated model, hard to train and prone to over-fitting. CNNs in contrast constrain a lot of weights to be identical and weight updates are applied jointly¹, thus giving a regularisation. Additionally, translation and limited rotation invariances are hard-wired into a CNN. CNNs were first used for handwritten digit recognition.

¹For implementation all weights are accessed via representative weights, this makes parameter data structures lean. Because of this CNNs are sometimes described to do *weight sharing*.

2.2 Network Architecture

2.2.1 Convolution

The first step is to imitate simple cells: in the hidden layer several *filters* or *feature maps* exist. The neurons of a feature map correspond to an area in the domain of the input image². All of these neurons only receive input from a local field of view due to sparse connectivity. The neurons' weights represent a linear shift-invariant filter that is applied to this field.

Calculating the weighted sum for all hidden neurons in a feature map is a discrete convolution. In 1 dimension, for a filter size (receptive field) of $2k + 1 = s$, the activation in the filter map is given by:

$$\mathbf{y}_j = h(\mathbf{a}_j) = h\left(\sum_{i=-k}^k x_i w_{j-i}\right) = h(\mathbf{x} * \mathbf{w})_j \quad (2.1)$$

Likewise for 2 dimensions \mathbf{w} is square matrix and a double-sum is needed for the convolution. If the input comes from a 2-dimensional feature map of size $s \times s$ with l different feature maps and the current layer exhibits m filters the weights of the convolution kernel are given by a tensor with shape $[m, l, s, s]$. It is worth noting that, unlike for a normal layer, the parameter number is not affected by the extent of the input at all. Moreover, the convolution is an operation of less computational cost compared to the matrix multiplication of a standard layer (provided the layer sizes are equal).




Figure 2.1: The flat “non-spatial” input neuron vector is reshaped into a matrix so that each neuron corresponds to a pixel in space. Within the neighbourhood of this pixel in the *previous* layer (illustrated in green for neuron x_7) the response of a filter (defined by connection weights) is computed. The response determines the neuron’s activation. As each neuron only has connections within a local window and the pattern of connection weights is the same for all neurons, this computation is equivalent to a convolution

The different feature maps can be imagined as a stack with a third dimension additional to the two spatial dimensions. Likewise do the filter kernels have a third dimension because

²Due to border effects the valid image region has some offset from the edges.

they can *combine* all the feature maps in the previous layer - but the convolution is still only across the two-dimensional domain.

As a consequence of weight sharing, all output neurons in a CNN process the input data of their corresponding receptive fields in exactly the same way. Therefore such a network can be viewed as many parallel networks with a single output neuron, glued next to each other on a grid. The receptive fields are overlapping (as they are bigger than the stride), hence creating one network is more efficient than many because the feature maps can be reused within the implementation.

There exist some variations of usage of convolutions: for input images with several channels (colour, depth etc.) the channels are processed like feature maps in the previous layer. Convolution kernels can also be even-sized, in which case a neuron in the feature map does not directly correspond to a pixel/neuron in the previous layer but “sits in between”. Convolutions need not be dense, meaning that it is possible to leave out neurons in the hidden layer, so that the receptive fields do not overlap; this is mainly done to reduce computation. Input can have higher dimensions than two. Then the convolution is required to be carried out across more dimensions, which increases computational cost.

2.2.2 Max-Pooling

The second step is inspired by the complex cells’ integration capability: if a filter detects a feature, such as a line with certain angle, it is of greater importance *what* is detected than *where* exactly it is. The fine spatial resolution of the filter maps is not needed and processing can be made more sparse and computationally efficient by reducing it. This also gives rise to translation invariance and slight rotational invariance. Note that this invariance is not to be confused with the translation invariance resulting from convolution: the latter ensures that the the same objects in an image will be detected the same way regardless of their position. The invariance due to pooling makes the combination of features invariant at the length scale of the pool size.

Plain sub-sampling is equivalent to a strided convolution. Alternatively the average across a pool in the feature map can be used. Comparisons by [LeCun and Bengio, 1995] have shown that taking the maximum of the pool outperforms taking the average and has therefore become the standard.

2.3 Operation of a CNN

For illustration refer to figure 2.3. Several convolutional and pooling layers with a variable number of filters and filter sizes are stacked together. In this part of the network the

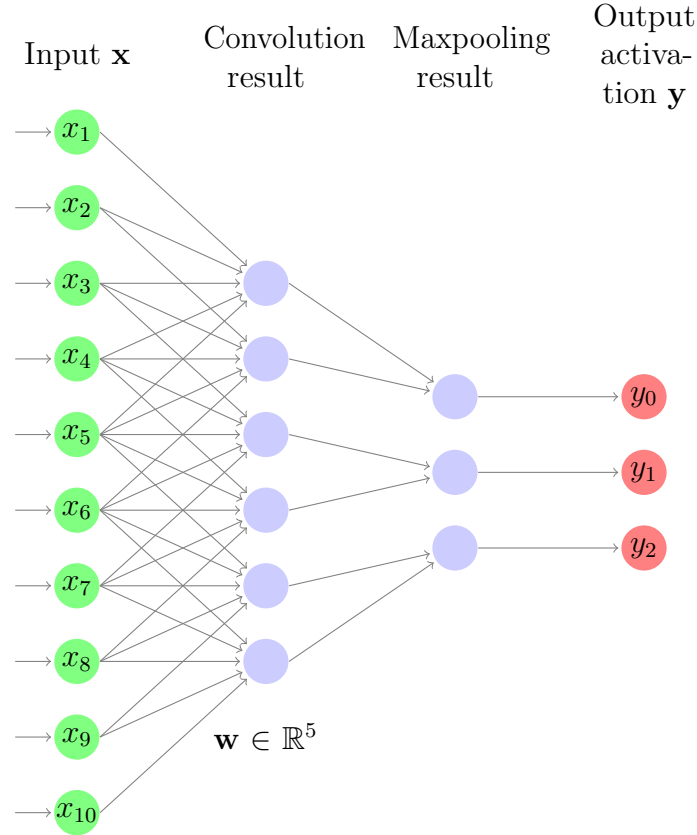


Figure 2.2: A one-dimensional convolutional and max-pooling layer (only one filter depicted): because all hidden units have the same weights, the input is convolved with a kernel (of size five). The max-pooling takes the maximum of a local pool (of size two here).

feature maps preserve the spatial correspondence to the input domain and each neuron can be identified by a position in space and as a member of a particular filter map. The pooling reduces the size of the filter maps exponentially and after a few layers there are relatively few neurons per layer.

This first part can be interpreted as a hierarchical feature extractor: simple edge-like filters of the first layers are combined to form more complex features (e.g. intersection of lines at certain angles, arcs etc.), from layer to layer the representations become more complex and due to pooling less located.

After the input has ideally been transformed to a sparse encoding by higher order abstract features this code can be used to train a classifier e.g. to give probabilities for integer values of handwritten digits. As we have already seen, normal NNs are universal classifiers so it is natural to use them: to the last convolutional layer a number of standard fully-connected layers are attached. This breaks the spatial structure and the separation into distinct filter maps as the all activations of the last convolutional layer are just flattened into a large

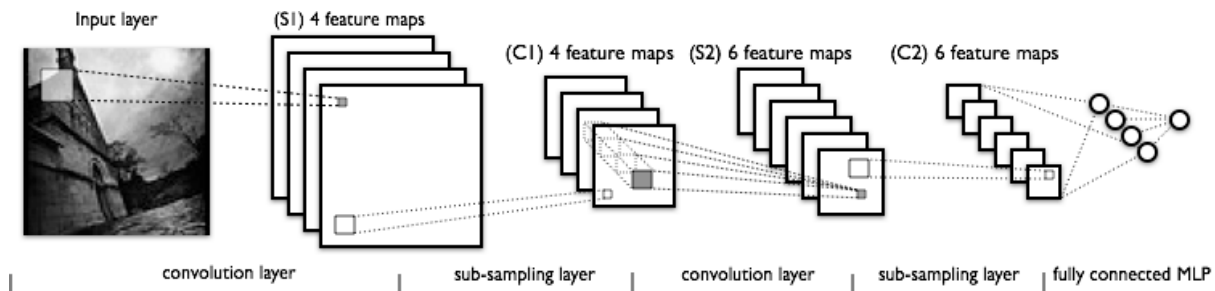


Figure 2.3: Schematic visualisation of a CNN. The “stacks” in each layer illustrate a number of feature maps. [source: <http://deeplearning.net/tutorial/lenet.html>]

vector.

The training procedure is in principle the same as for normal NNs. Like for normal NNs the initialisation of the convolutional filters is usually random and for quite often Gabor-like filters automatically evolve during optimisation. In other cases Gabor-Filters can be used as initial values or filters that have separately been learnt e.g. by k-means or PCA.

2.3.1 Providing training data to a CNN for image-to-image tasks

Patch creation

Generally, feeding the whole image to a NN is computationally intractable with current hardware, and it also creates problems for images of different sizes. Instead the image is cut into smaller, square *patches* of constant size. Depending on the patch size, multiple patches can be stacked to form a mini-batch which is used for Backprop-training. Using all patches from all training images as in full-batch training is also computationally infeasible.

Many objects in images are invariant under congruence transformations i.e. translation, reflection and rotation. This can be exploited by *augmenting* the available data by application of these transformations. When training the the cutting of patches is subject to a random offsets in the whole image, thus yielding translated patches. Reflections were applied by flipping axes of the images. This can be done in constant time by creating strided memviews of the images. Rotations³ and other deformations are computationally more intense and can be done in background processes.

³except by 90°

Label creation

A CNN with no fully-connected layers at the end has multiple output neurons, corresponding to a grid in image space. Due to the max-pooling, neurons in the output layer of a deep CNN are separated by more than one pixel in the original image space. If neurons are aligned so that they lie at the center pixel of their receptive fields, the pixel distance of two output neurons is called *stride*; e.g. after three times max-pooling by a factor of two the resulting stride is $2^3 = 8$. The labels to which the output neurons are trained are determined by sub-sampling the labels with the corresponding stride (see figure 2.4).

Independent of the pooling, the total receptive field is even bigger because the convolutional filters take into account neighbouring neurons/pixels in each layer, as a side effect this causes a boundary region in the input domain for which no predictions can be made⁴.

To obtain a dense prediction for each pixel in the image, not just a strided one, it is necessary to feed the input patch several times through the network and shift it by one pixel. A dense image is generated by combining predictions of the output neurons. As this is computationally inefficient max-fragment-pooling (MFP) can be used which calculates all possible shifts in one pass.

⁴To make a prediction there, a patch must be created s.t. this region does not lie on the boundary. Alternatively good results can be achieved by mirroring the boundary.

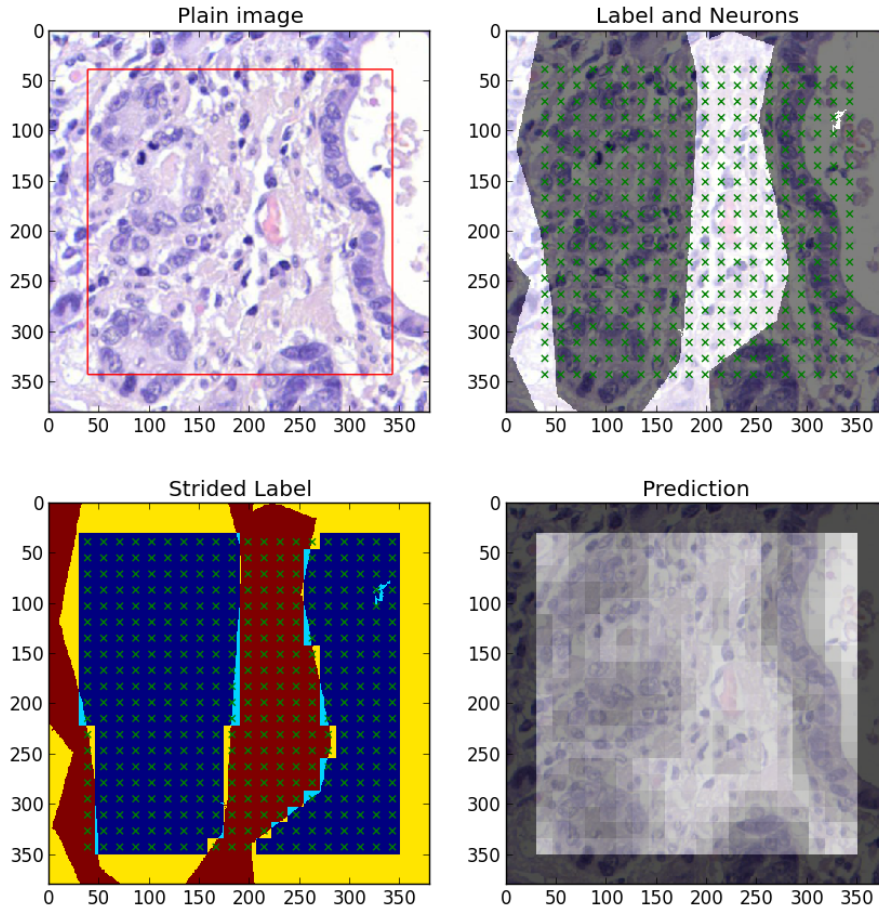


Figure 2.4: CNN data for image-to-image tasks:

top left Input image with border region indicated: due to convolution borders a stripe occurs for which no predictions are made.

top right Image with true label overlaid. In green is the strided array of output neurons for a particular CNN indicated.

bottom left From the true labels, training labels were created for each neuron by subsampling the labels (For better visualisation the labels were scaled up to cover the whole block as opposed to only the center pixel)

bottom right Strided predictive probability map from a trained NN. Each “block” corresponds to the output of one neuron at its center. (For better visualisation the predictions were scaled up to cover the whole block as opposed to only the center pixel).

List of Figures

1.1	A single neuron	5
1.2	Simple 2-layer network	7
1.3	Comparison of error-functions	11
1.4	Gradient Descent	13
2.1	Reshaping networks from 1d to 2d	18
2.2	Convolution and Max-Pooling Layer	20
2.3	Convolutional Neural Network	21
2.4	CNN data for image-to-image tasks	23

Bibliography

- [Bishop et al., 2006] Bishop, C. M. et al. (2006). *Pattern recognition and machine learning*, volume 1. springer New York.
- [Cybenko, 1989] Cybenko, G. (1989). Approximation by superpositions of a sigmoidal function. *Mathematics of control, signals and systems*, 2(4):303–314.
- [Fukushima, 1980] Fukushima, K. (1980). Neocognitron: A self-organizing neural network model for a mechanism of pattern recognition unaffected by shift in position. *Biological cybernetics*, 36(4):193–202.
- [Hodgkin and Huxley, 1952] Hodgkin, A. L. and Huxley, A. F. (1952). A quantitative description of membrane current and its application to conduction and excitation in nerve. *The Journal of physiology*, 117(4):500.
- [Hubel and Wiesel, 1962] Hubel, D. H. and Wiesel, T. N. (1962). Receptive fields, binocular interaction and functional architecture in the cat’s visual cortex. *The Journal of physiology*, 160(1):106.
- [Lazar et al., 2009] Lazar, A., Pipa, G., and Triesch, J. (2009). Sorn: a self-organizing recurrent neural network. *Frontiers in computational neuroscience*, 3.
- [LeCun and Bengio, 1995] LeCun, Y. and Bengio, Y. (1995). Convolutional networks for images, speech, and time series. *The handbook of brain theory and neural networks*, 3361.
- [LeCun et al., 1998] LeCun, Y., Bottou, L., Bengio, Y., and Haffner, P. (1998). Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11):2278–2324.
- [LeCun et al., 2012] LeCun, Y. A., Bottou, L., Orr, G. B., and Müller, K.-R. (2012). Efficient backprop. In *Neural networks: Tricks of the trade*, pages 9–48. Springer.
- [McCulloch and Pitts, 1943] McCulloch, W. S. and Pitts, W. (1943). A logical calculus of the ideas immanent in nervous activity. *The bulletin of mathematical biophysics*, 5(4):115–133.

- [Rosenblatt, 1958] Rosenblatt, F. (1958). The perceptron: a probabilistic model for information storage and organization in the brain. *Psychological review*, 65(6):386.
- [Rumelhart et al., 1986] Rumelhart, D. E., Hinton, G. E., , and Williams, R. J. (1986). Learning representations by back-propagating errors. *Nature*, 323:533–536.
- [Young, 1938] Young, J. Z. (1938). The functioning of the giant nerve fibres of the squid. *Journal of Experimental Biology*, 15(2):170–185.